

National University of Singapore
School of Computing
CS3243 Introduction to AI

Tutorial 6: Adversarial Search


Issued: Week 8

Discussion in: Week 9

Important Instructions:

- **Tutorial Assignment 6** consists of **Question 1** from this tutorial.
- Your solution(s) must be TYPE-WRITTEN, though diagrams may be hand-drawn.
- You must submit your solution(s) via **Canvas > Assignments**, satisfying the deadlines:
 - **Pre-tutorial 6 submission by Week 8 Sunday, 2359 hrs.**
 - **Post-tutorial 6 submission by Week 9 Friday, 2359 hrs.**
- You must make both submissions for your assignment score to be counted.


1. Consider the **Minimax** tree given in the figure below. In the figure, black nodes are controlled by the MAX player, and white nodes are controlled by the MIN player. Payments (terminal nodes) are squares, where the number within denotes the amount that the MIN player pays to the MAX player (an amount of 0 means that MIN pays nothing to MAX). Naturally, MAX wants to maximise the amount they receive, and MIN wants to minimise the amount they pay.



Suppose that we use the **α - β Pruning** algorithm, given in Figure 5.7 of AIMA 4th edition (reproduced in the **Appendix**).



- (a) *Assume that we iterate over nodes from right to left*; mark with an “X” all edges that are pruned by α - β pruning, if any.
- (b) *Assume that we iterate over nodes from left to right*; mark with an “X” all edges that are pruned by α - β pruning, if any.
- (c) Determine if the effectiveness of pruning depends on iteration order (i.e., provide a **yes or no** answer. (There is no need to elaborate on this answer for your assignment submission. However, you should consider why iteration order is important for discussion during the tutorial.)

2. Assume that α - β (alpha-beta) pruning is applied to each of the following game trees in a left-to-right exploration. Further, assume zero-sum games. Select the options where the given pruning is **possible** under some choice of pay-off values. Note that the crossed-off edges indicate where pruning has occurred.



Option G: None of the above.

3. A game tree is **best-score-first** ordered if for all non-terminal nodes, its first successor represents the best possible move for that position. The tree on the right is an example of a game tree that is **best-score-first** ordered. Notice that at the first node (which is a MAX node), the first action is best as it leads to a utility score of 7 while the second action leads to a utility score of 3. In the second node (which is a MIN node), the first action is best as it leads to a utility score of 3 while the second action leads to a utility score of 6. In the third node (which is a MAX node), the first action is once again best. Hence, the above tree is **best-score-first** ordered.



Note that a different agent can use a different move-ordering heuristic to generate a different game tree for the same game. For example, the game tree on the left might be generated instead. This is a permutation of the first tree. Note that there are many possible permutations for a given game tree, all of which corresponds to a slightly different move-ordering heuristic being used. Notice how in the first tree (which is **best-score-first** ordered), one branch is pruned by the α - β Pruning algorithm. However, in the second tree (which is a permutation, but not **best-score-first** ordered), no branch is pruned by the α - β Pruning algorithm. Hence, fewer nodes are visited with the first tree than the second tree.

Prove or disprove the following statement.

Let T be a **best-score-first** ordered game tree. Suppose T' is a permutation of T such that it is not **best-score-first** ordered. If the α - β Pruning algorithm visits N nodes given T , then it will visit $M \geq N$ nodes given T' .

4. With the **Minimax** algorithm, we know that the value v , computed at the root (i.e., the utility for the MAX player), is a worst-case value. This means that if the opponent MIN does not act optimally, the actual outcome v' for MAX can only be better, and never worse than v . That said, the **Minimax** algorithm may not select the optimal move given sub-optimal play from the MIN player.

Construct an example where, should the MIN player play sub-optimally, the **Minimax** algorithm makes a sub-optimal move.

Appendix

The following is a description of the **α - β Pruning** algorithm. It is referenced from the AIMA 4th edition textbook (Figure 5.7).

```

function ALPHA-BETA-SEARCH(game, state) returns an action
  player  $\leftarrow$  game.TO-MOVE(state)
  value, move  $\leftarrow$  MAX-VALUE(game, state,  $-\infty$ ,  $+\infty$ )
  return move

function MAX-VALUE(game, state,  $\alpha$ ,  $\beta$ ) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v  $\leftarrow$   $-\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MIN-VALUE(game, game.RESULT(state, a),  $\alpha$ ,  $\beta$ )
    if v2 > v then
      v, move  $\leftarrow$  v2, a
       $\alpha \leftarrow$  MAX( $\alpha$ , v)
    if v  $\geq$   $\beta$  then return v, move
  return v, move

function MIN-VALUE(game, state,  $\alpha$ ,  $\beta$ ) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v  $\leftarrow$   $+\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MAX-VALUE(game, game.RESULT(state, a),  $\alpha$ ,  $\beta$ )
    if v2 < v then
      v, move  $\leftarrow$  v2, a
       $\beta \leftarrow$  MIN( $\beta$ , v)
    if v  $\leq$   $\alpha$  then return v, move
  return v, move

```